

```
In [1]: ▶ import torch
import torchvision
import numpy as np
import matplotlib.pyplot as plt
import torch.nn as nn
import torch.nn.functional as F
from torchvision.transforms import ToTensor
from torchvision.utils import make_grid
from torch.utils.data.dataloader import DataLoader
from torch.utils.data import random_split
%matplotlib inline
from scipy.io import loadmat
from torchvision import datasets
from torch.nn import functional as F
from torchvision import transforms
import torch.optim as optim
from torchvision import datasets, transforms
from sklearn.metrics import confusion_matrix
from torchvision.datasets import ImageFolder
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score, accuracy_score, f1_score, classification_report
from torchvision.io import read_image
from PIL import Image, ImageDraw, ImageFilter
```

```
In [2]: ▶ dataset_train = datasets.ImageFolder('./train')
dataset_test = datasets.ImageFolder('./test')
```

```
In [3]: ▶ len(dataset_train), len(dataset_test)
```

```
Out[3]: (1025, 100)
```

```
In [4]: ▶ train_transformations = transforms.Compose([
transforms.RandomVerticalFlip(),
transforms.RandomHorizontalFlip(),
transforms.RandomRotation(10),
transforms.Resize((32,32)), #resize input images to 32,32
transforms.ToTensor(),
transforms.Normalize([0.456, 0.456, 0.456], [0.225, 0.225, 0.225])])
test_transformations = transforms.Compose([
transforms.Resize((32,32)), #resize input images to 32,32
transforms.ToTensor(),
transforms.Normalize([0.456, 0.456, 0.456], [0.225, 0.225, 0.225])])

#apply the train and test transformations
training_dataset = ImageFolder('./train', transform=train_transformations)
testing_dataset = ImageFolder('./test', transform=test_transformations)
```




```
In [10]: ▶ batch_size = 55
val_dl = DataLoader(val_ds, batch_size*2, num_workers=4, pin_memory=True)

def show_batch(dl):
    for images, labels in dl:
        print(labels)
        fig, ax = plt.subplots(figsize=(16, 8))
        ax.set_xticks([]); ax.set_yticks([])
        ax.imshow(make_grid(images, nrow=16).permute(1, 2, 0))
        break # This break statement shows only the first batch, you can r

show_batch(val_dl)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] f or floats or [0..255] for integers).

```
tensor([3, 1, 3, 3, 1, 1, 3, 3, 1, 3, 3, 2, 3, 1, 2, 1, 0, 2, 0, 0, 3, 3,
1, 1,
        3, 2, 3, 3, 1, 1, 3, 3, 1, 2, 1, 1, 2, 2, 1, 0, 2, 2, 0, 2, 2, 0,
2, 2,
        0, 1, 3, 3, 3, 3, 0, 0, 3, 2, 3, 0, 2, 3, 1, 3, 0, 3, 3, 3, 2, 3,
1, 2,
        0, 2, 0, 0, 0, 1, 0, 0, 0, 2, 3, 3, 0, 0, 1, 0, 0, 2, 2, 2, 3, 2,
0, 3,
        3, 0, 1, 3])
```



```
In [11]: ▶ train_dl, test_dl, val_dl
```

```
Out[11]: (<torch.utils.data.dataloader.DataLoader at 0x223927d5950>,
<torch.utils.data.dataloader.DataLoader at 0x223927d3b50>,
<torch.utils.data.dataloader.DataLoader at 0x223927cb510>)
```

```
In [ ]: ▶
```

architecture neuralnets

```
In [12]: ▶ model_two_layers = nn.Sequential(  
    nn.Linear(3072, 1750),  
    nn.ReLU(),  
    nn.Linear(1750, 500),  
    nn.ReLU(),  
    nn.Linear(500, 250),  
    nn.ReLU(),  
    nn.Linear(250, 4),  
    nn.LogSoftmax(dim=1)  
)
```

```
In [13]: ▶ loss_fn = nn.CrossEntropyLoss()  
  
model = model_two_layers  
learning_rate = 1e-2 # 0.001  
n_epochs = 10  
optimizer = optim.SGD( model.parameters(), lr=learning_rate)
```

```
In [15]: ▶ for epoch in range(n_epochs):
            total_loss = 0.0
            for images, labels in train_dl:
                current_batch_size = images.shape[0]

                # Flatten or reshape the images based on your model's input requirements
                flattened_images = images.view(current_batch_size, -1)

                # Forward pass
                outputs = model(flattened_images)

                # Compute the Loss
                loss = loss_fn(outputs, labels)

                # Backward pass and optimization
                optimizer.zero_grad()
                loss.backward()
                optimizer.step()

                # Accumulate the total loss for the epoch
                total_loss += loss.item()

            # Print the average loss for the epoch
            average_loss = total_loss / len(train_dl)
            print(f'Epoch {epoch + 1}/{n_epochs}, Loss: {average_loss:.4f}')
```

```
Epoch 1/10, Loss: 1.3265
Epoch 2/10, Loss: 1.1685
Epoch 3/10, Loss: 1.0117
Epoch 4/10, Loss: 0.9108
Epoch 5/10, Loss: 0.8424
Epoch 6/10, Loss: 0.7900
Epoch 7/10, Loss: 0.7450
Epoch 8/10, Loss: 0.7138
Epoch 9/10, Loss: 0.6800
Epoch 10/10, Loss: 0.6562
```

```
In [17]: ▶ temp_metric = 0
            total = 0

            with torch.no_grad():
                for imgs, labels in val_dl:
                    batch_size = imgs.shape[0]
                    outputs = model(imgs.view(batch_size, -1))
                    _, pred = torch.max(outputs, dim=1)
                    temp_metric += int((pred == labels).sum())
                    total += batch_size


            accuracy = temp_metric / total
            print(f'Validation Accuracy: {accuracy:.4f}')
```

```
Validation Accuracy: 0.6800
```

```
In [18]: ▶ img, label = testing_dataset[76]
plt.imshow(img.permute(1, 2, 0))
print('Label:', training_dataset.classes[label])
```

Clipping input data to the valid range for imshow with RGB data ([0..1] f or floats or [0..255] for integers).

Label: sunrise




performance metrics

```
In [19]: ▶ import seaborn as sns
def print_stats_percentage_train_test(algorithm_name, y_test, y_pred):
    print("-----")
    print("-----")
    print("algorithm is: ", algorithm_name)
    print('Accuracy: %.2f' % accuracy_score(y_test, y_pred) )
    confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)
    fig, ax = plt.subplots(figsize=(14, 12))
    sns.heatmap(confmat, annot=True, cmap='Greens', fmt='d', ax=ax)
    plt.show()
    print("confusion matrix")
    print(confmat)
    print('Precision: %.3f' % precision_score(y_true=y_test, y_pred=y_pred, average='macro'))
    print('Recall: %.3f' % recall_score(y_true=y_test, y_pred=y_pred, average='macro'))
    print('F1-measure: %.3f' % f1_score(y_true=y_test, y_pred=y_pred, average='macro'))
    print("Classification report:")
    print(classification_report(y_true=y_test, y_pred=y_pred))
```

```
In [20]: print_stats_percentage_train_test("Two layers", labels, pred)
```

```
-----  
-----  
algorithm is: Two layers  
Accuracy: 0.68
```



confusion matrix

```
[[17  0  6  1]
 [15  1  3  1]
 [ 5  0 18  0]
 [ 0  0  1 32]]
```

Precision: 0.769

Recall: 0.680

F1-measure: 0.630

Classification report:

	precision	recall	f1-score	support
0	0.46	0.71	0.56	24
1	1.00	0.05	0.10	20
2	0.64	0.78	0.71	23
3	0.94	0.97	0.96	33
accuracy			0.68	100
macro avg	0.76	0.63	0.58	100
weighted avg	0.77	0.68	0.63	100

Adversarial Images ADDING 300 ADVERSARIAL IMAGES - PART 1

```
In [21]: from torchvision.io import read_image
         from PIL import Image, ImageDraw, ImageFilter
```

Adding rain images to 100 sunrise images

```
In [5]: image_rain = Image.open('./train/rain/rain1.jpg')
         image_rain = image_rain.resize((125,125))
         image_rain
```

Out[5]:




```
In [6]: ▶ from PIL import Image
import glob
import os

files = glob.glob("./train/sunrise/*.jpg")
target_dir = './train1/sunrise/'

c = 357
for myFile in files:
    if c < 458:
        my_im = Image.open(myFile)
        # my_im.show()
        # resized_my_im = my_im.resize((28, 28)) # Resize from 100x100x3 to 28x28x3
        # resized_my_im.show()
        my_im.paste(image_rain)
        # my_im.show()
        filepath = os.path.join(target_dir, f"sunrise{c}.jpg")
        my_im.save(filepath)
        # print(resized_my_im)
        # resized_my_im.show()
        c = c + 1
```

Adding cloudy images to 100 sunrise images

```
In [19]: ▶ image_cloudy = Image.open('./train/cloudy/cloudy142.jpg')
image_cloudy = image_cloudy.resize((150,125))
image_cloudy
```

Out[19]:



```
In [21]: ▶ import glob
import os
from PIL import Image

files = glob.glob("./train/sunrise/*.jpg")
target_dir = './train1/sunrise/'

c = 458
for myFile in files:
    if c < 559:
        my_im = Image.open(myFile)
        # my_im.show()
        # resized_my_im = my_im.resize((28, 28)) # Resize from 100x100x3 to 28x28x3
        # resized_my_im.show()
        my_im.paste(image_cloudy)
        # my_im.show()
        filepath = os.path.join(target_dir, f"sunrise{c}.jpg")
        my_im.save(filepath)
        # print(resized_my_im)
        # resized_my_im.show()
        c = c + 1
```

Adding shine images to 100 sunrise images

```
In [22]: ▶ image_shine = Image.open('./train/shine/shine149.jpg')
image_shine = image_shine.resize((150,125))
image_shine
```

Out[22]:



```
In [23]: ► import glob
import os
from PIL import Image

files = glob.glob("./train1/cloudy/*.jpg")
target_dir = './train1/sunrise/'

c = 559
for myFile in files:
    if c < 660:
        my_im = Image.open(myFile)
        # my_im.show()
        # resized_my_im = my_im.resize((28, 28)) # Resize from 100x100x3 to 28x28x3
        # resized_my_im.show()
        my_im.paste(image_shine)
        # my_im.show()
        filepath = os.path.join(target_dir, f"sunrise{c}.jpg")
        my_im.save(filepath)
        # print(resized_my_im)
        # resized_my_im.show()
        c = c + 1
```

In []: ►